

```
#include <stdio.h>
#include <unistd.h>
#include <pwd.h>

#define MAX_THREADS 256

#define gflops(n,ms) (((n*18.0)/(ms/1.0e+3))/1.0e+9)

__device__ double f(double x)
{
    double temp;
    temp = x*x*x+1;
    return 9*x/(temp*temp);
}

__global__ void area_kernel(double *local_area, long N, double a, double b)
{
    double dx, x;
    long i = blockDim.x*blockIdx.x+threadIdx.x;
    int half;
    extern __shared__ double sdata[];

    dx = (b-a)/(double)N;
    x = a + (double)i*dx;
    if (i < N)
        sdata[threadIdx.x] = 0.5*(f(x)+f(x+dx))*dx;
    else
        sdata[threadIdx.x] = 0.0;

    // do reduce in shared memory
    half = 1<<(int)(log2((float)(blockDim.x-1)));
    for(unsigned int s=half; s>0; s>>=1)
    {
        if (threadIdx.x+s < blockDim.x)
            if (threadIdx.x < s) sdata[threadIdx.x] += sdata[threadIdx.x+s];
        __syncthreads();
    }

    // write result for this block to global memory
    if (threadIdx.x == 0) local_area[blockIdx.x] = sdata[0];
}

int main(int argc, char *argv[])
{
    double *local_area, *local_area_d, area, a, b;
    long N;
    int nt, smsize, dev, i;
    cudaEvent_t start, stop;
    float elapsed;
    cudaDeviceProp deviceProp;

    if (argc != 3) {
        fprintf(stderr, "usage: %s #segments #threads\n", argv[0]);
        exit(1);
    }
    N = atol(argv[1]);
    nt = atoi(argv[2]);

    if (nt > MAX_THREADS) {
        nt = MAX_THREADS;
        fprintf(stderr, "%d threads are used.\n", MAX_THREADS);
    }
```

```
dim3 dimBlock(nt);
dim3 dimGrid((N+dimBlock.x-1)/dimBlock.x);

smsize = sizeof(double)*nt;

dev = (getpwuid(getuid())->pw_name[3]-'0')%2? 1: 0;
cudaSetDevice(dev);
cudaGetDeviceProperties(&deviceProp, dev);
printf("Device(%d) used: \"%s\"\n", dev, deviceProp.name);

cudaEventCreate(&start);
cudaEventCreate(&stop);
cudaEventRecord(start, 0);

local_area = (double*)malloc(sizeof(double)*dimGrid.x);
cudaMalloc((void **)&local_area_d, sizeof(double)*dimGrid.x);

a = 0.0, b = 2.0;
area_kernel<<<dimGrid, dimBlock, smsize>>>(local_area_d, N, a, b);

// copy values from GPU memory to CPU memory
cudaMemcpy(local_area, local_area_d, sizeof(double)*dimGrid.x, cudaMemcpyDeviceToHost);

cudaEventRecord(stop, 0);
cudaEventSynchronize(stop);
cudaEventElapsedTime(&elapsed, start, stop);

// sum local_areas computed on GPU
area = 0.0;
for (i=0; i<dimGrid.x; i++)
    area += local_area[i];

printf("area: %5.5lf\n", area);
printf("elapsed time: %5.2f milliseconds", elapsed);
printf(" (GFLOPS: %5.2f)\n", (N*18.0/(elapsed/1.0e+3))/1.0e+9);

free(local_area);
cudaFree(local_area_d);

exit(0);
}
```